

Yi Hou (厚熠)

Beijing, China · houyi25@mails.ucas.ac.cn · (+86) 181-9337-4598 · github.com/Helios-YQH

Second-year Artificial Intelligence undergraduate working across the full LLM stack: Triton kernels, distributed training, pretraining data pipelines, and post-training. Author of five public technical reports built on controlled experiments.

EDUCATION

University of Chinese Academy of Sciences (UCAS) — Beijing, China

B.Eng. in Artificial Intelligence · Sep 2025 – Jun 2029 (expected) · CET-6 (College English Test Band 6) passed

SKILLS

- **Languages:** Python (primary), C/C++, JavaScript/TypeScript, LaTeX
- **LLM & Training:** PyTorch, Jittor; Transformer architectures from scratch; Triton kernels (FlashAttention-2); SFT / GRPO / DPO; RAG
- **Systems:** DDP / ZeRO-1 / FSDP implemented from scratch; mixed precision (bf16/fp16); Nsight Systems profiling
- **Engineering:** FastAPI, SQLite, Chroma, Git, Linux

PROJECTS

LLM from Scratch — 4 Projects · 5 Technical Reports (Stanford CS336, self-study)

Public monorepo: github.com/Helios-YQH/llm-from-scratch

- **Systems:** 5.7× faster attention than naive attention at 64K context (bf16) with hand-written Triton FlashAttention-2 kernels — fp32 64K contexts run where naive attention OOMs; from-scratch distributed stack: DDP variants (communication overlap cut the gradient-communication share from 69% to 59% on 4 GPUs), ZeRO-1 (~23.7% memory), FSDP (~43% steady-state memory, cross-checked against PyTorch's official FSDP); Nsight profiling across 128M–3.4B-parameter models (bf16 mixed precision ~2–3×)
- **Pretraining:** trained a 22.7M-parameter decoder-only Transformer to 1.57 val loss in 1.5 GPU-hours on a 6×A6000 cluster at ICT CAS; bit-exact BPE tokenizer (4.18 bytes/token on TinyStories); controlled ablations over normalization / positional encoding / activation identified that a cosine LR tail can mask divergence
- **Data:** quantified crawled-corpus filtering with a paired ablation — -0.27 nats validation loss (perplexity 282 vs 367) — over a pipeline turning 2.14M Common Crawl documents into 340M GPT-2 tokens (HTML extraction, language ID, PII masking, quality classifiers, exact + MinHash-LSH dedup)
- **GRPO:** controlled sweeps on OLMo-2-1B / GSM8K establish that prompt design — not the advantage estimator — is the first-order variable at 1B scale, and that RL retains ~90–105% of its gain under a 10%-flip verifier where test-time selection collapses to 57%
- **SFT + DPO:** quantified judge dependence in preference evaluation — AlpacaEval win rate for the same LoRA-SFT + DPO Llama-3.1-8B swings from 3.8% to 33.2% across four judges; measured MMLU / GSM8K capability trade-offs and SimpleSafetyTests safety

ZhiyanNavigator — LLM-Powered Paper Learning Assistant

Backend developer · 2026 China International College Students' Innovation Competition

- Built the backend learning loop of a paper-reading assistant: paper ingestion (arXiv link/ID or local PDF, with download and caching) → LLM-generated proficiency assessment → tiered study roadmap → card-by-card explanations with self-check quizzes
- Implemented RAG (local bge-small-zh embeddings + Chroma retrieval) and SSE-streaming LLM endpoints on FastAPI, which also serves the web UI; refactored an initially over-complex workflow into a minimal, reliable end-to-end pipeline

Diffusion Models from Scratch — UNet vs. DiT, ϵ -Prediction vs. Rectified Flow (PyTorch)

Public repo: github.com/Helios-YQH/diffusion-from-scratch

- Rectified flow beats ϵ -prediction at every budget in a controlled UNet-vs-DiT comparison (best MNIST FID 0.88); DDIM at 50 steps matches the 1,000-step ancestral FID within 0.8% at ~20× the speed; negative results documented (DDIM degrades DiT); multi-GPU DDP over 5 GPUs, and a 174M-param CelebA UNet at FID 11.4

Point Cloud Denoising — StraightPCF (CVPR 2024) Reproduction in Jittor

Solo · 6th Jittor AI Challenge, Apr–Jul 2026 — national top 100 · github.com/Helios-YQH/jittor-competition-point-cloud

- Reproduced StraightPCF in Jittor — dynamic EdgeConv encoder, coupled velocity modules, and distance-modulated Euler integration (~530K parameters) predicting per-point displacements back onto the surface; trained on ~20K ShapeNet meshes, cutting Chamfer distance by 51% and point-to-surface error by 67% versus the noisy input (best competition score 67.44)

Additional Projects

- **CUMCM 2026 (national mathematical-modeling contest, Problem A):** modeled 2D axisymmetric coupled heat-and-mass transfer for herb drying with implicit finite-difference solvers and nonlinear, state-dependent material properties; completed the modeling, numerical implementation, and LaTeX paper
- **3D Gaussian Splatting:** COLMAP reconstruction and 3DGS training pipeline